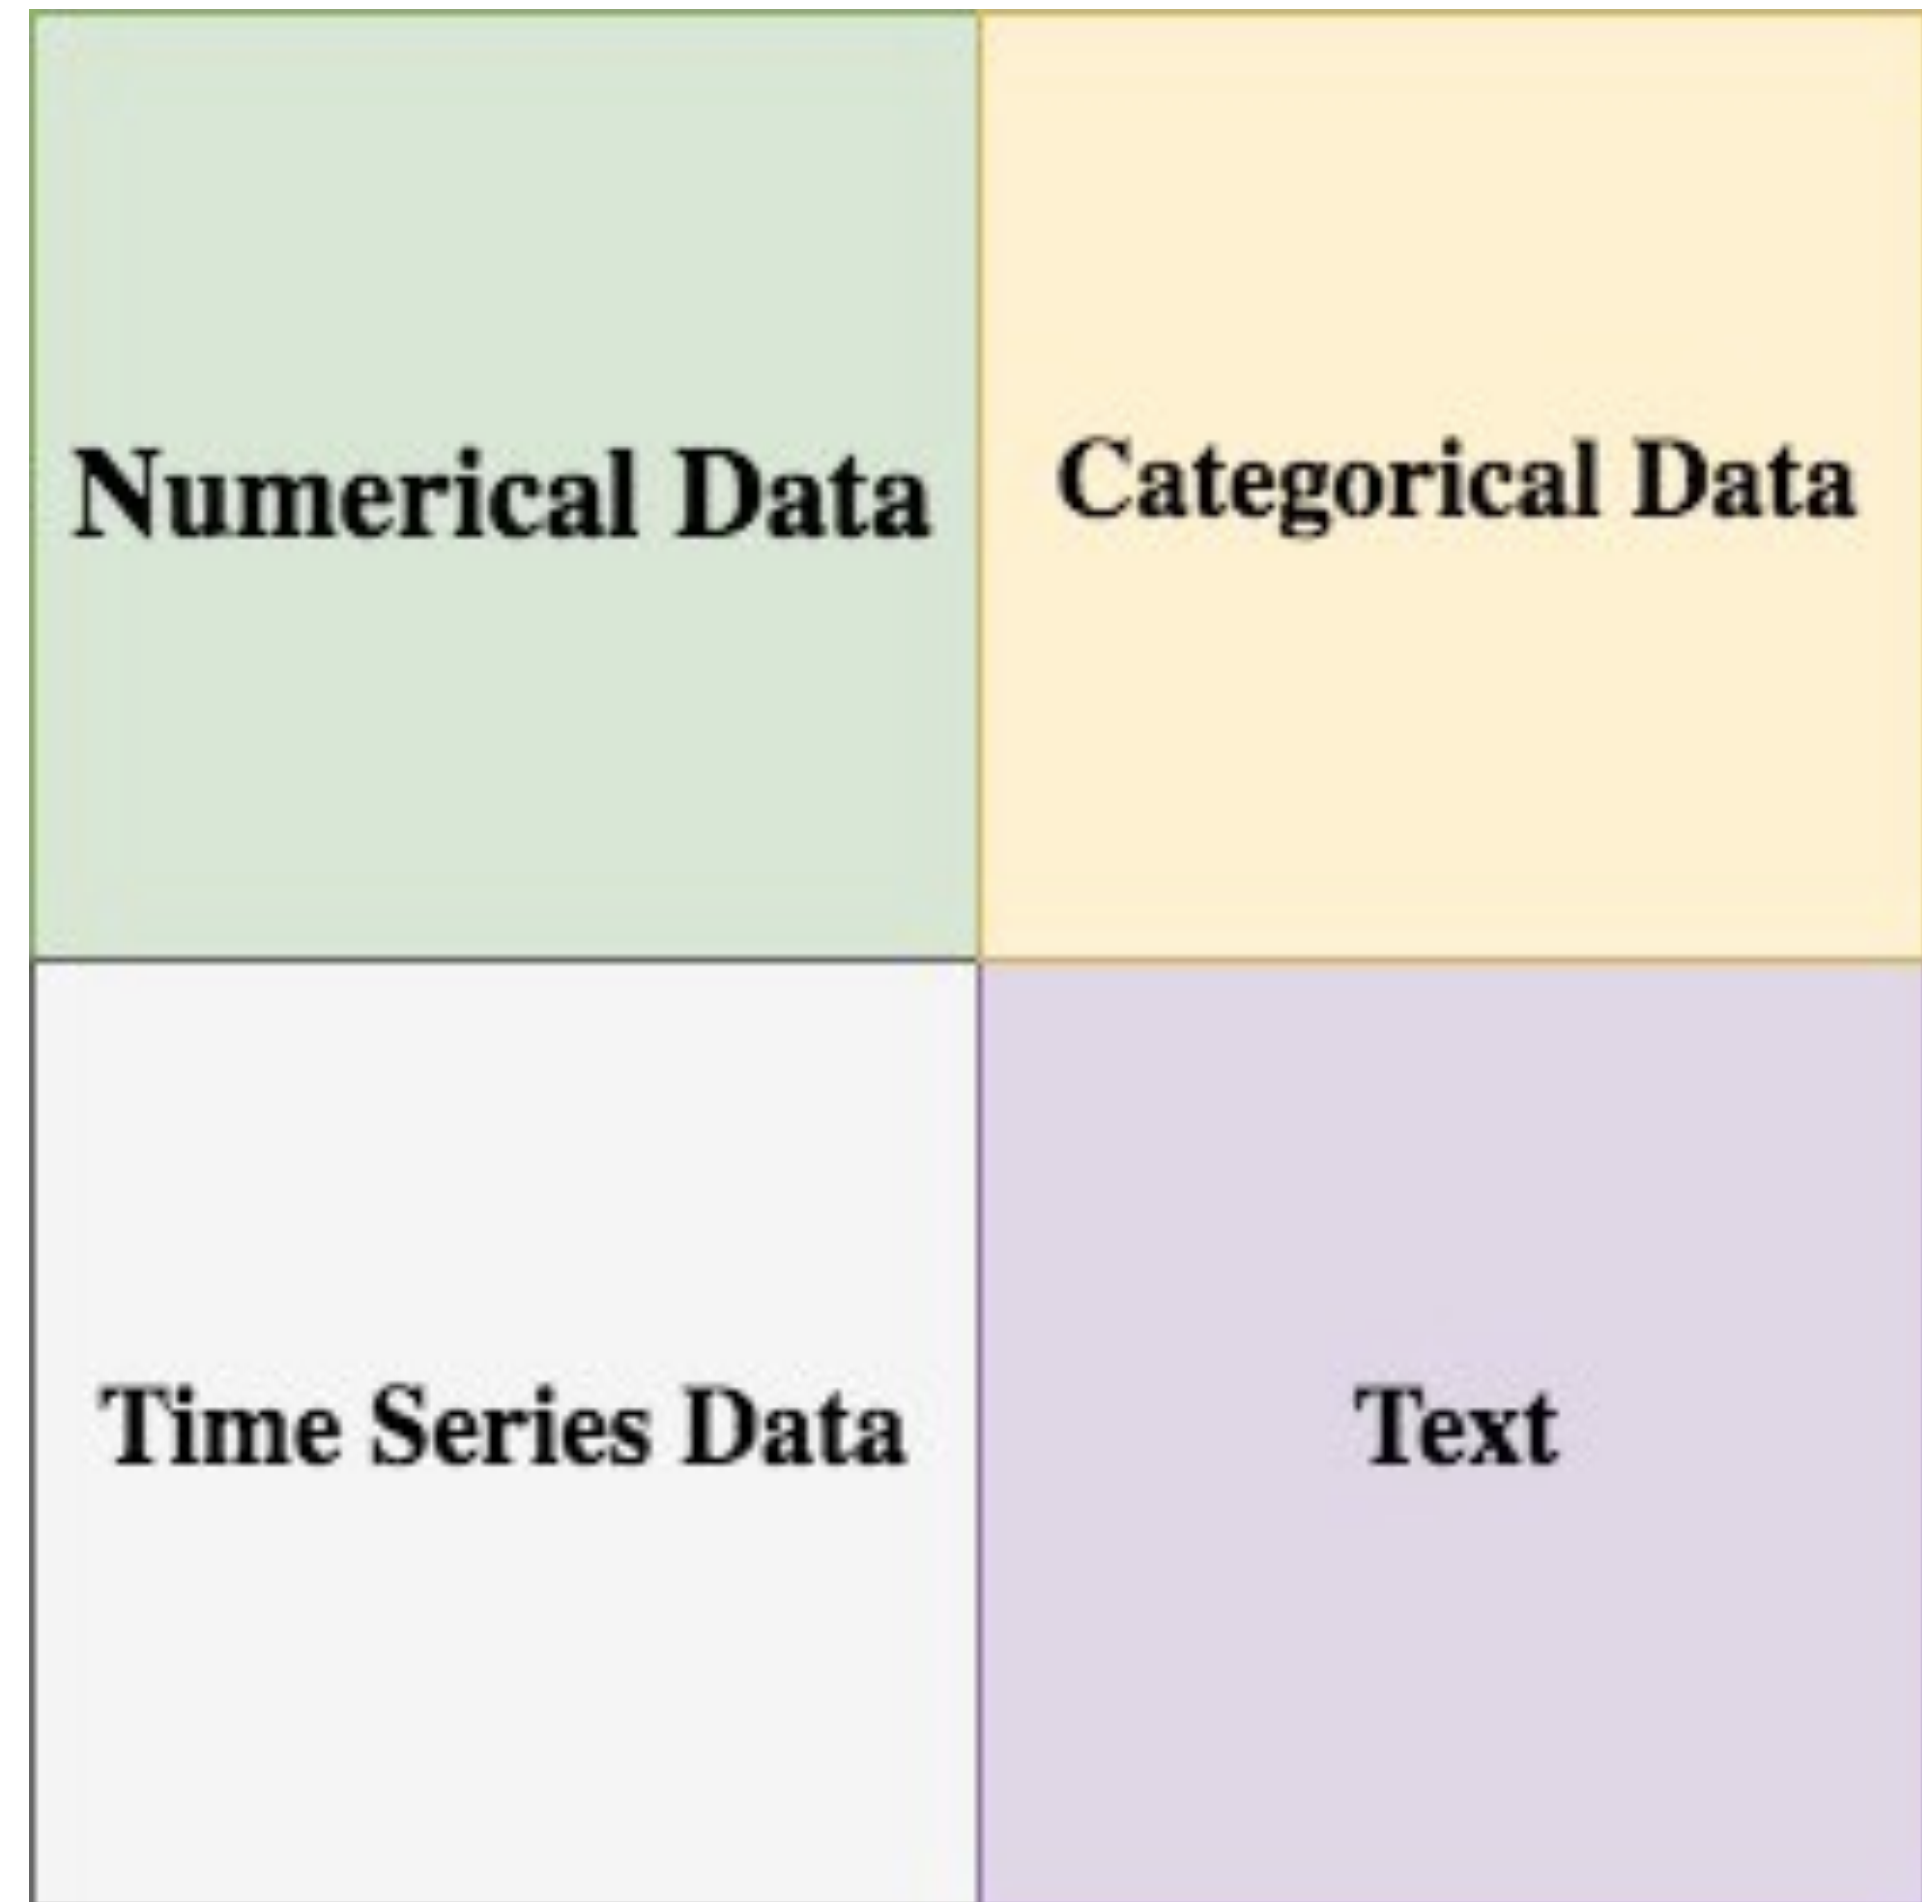


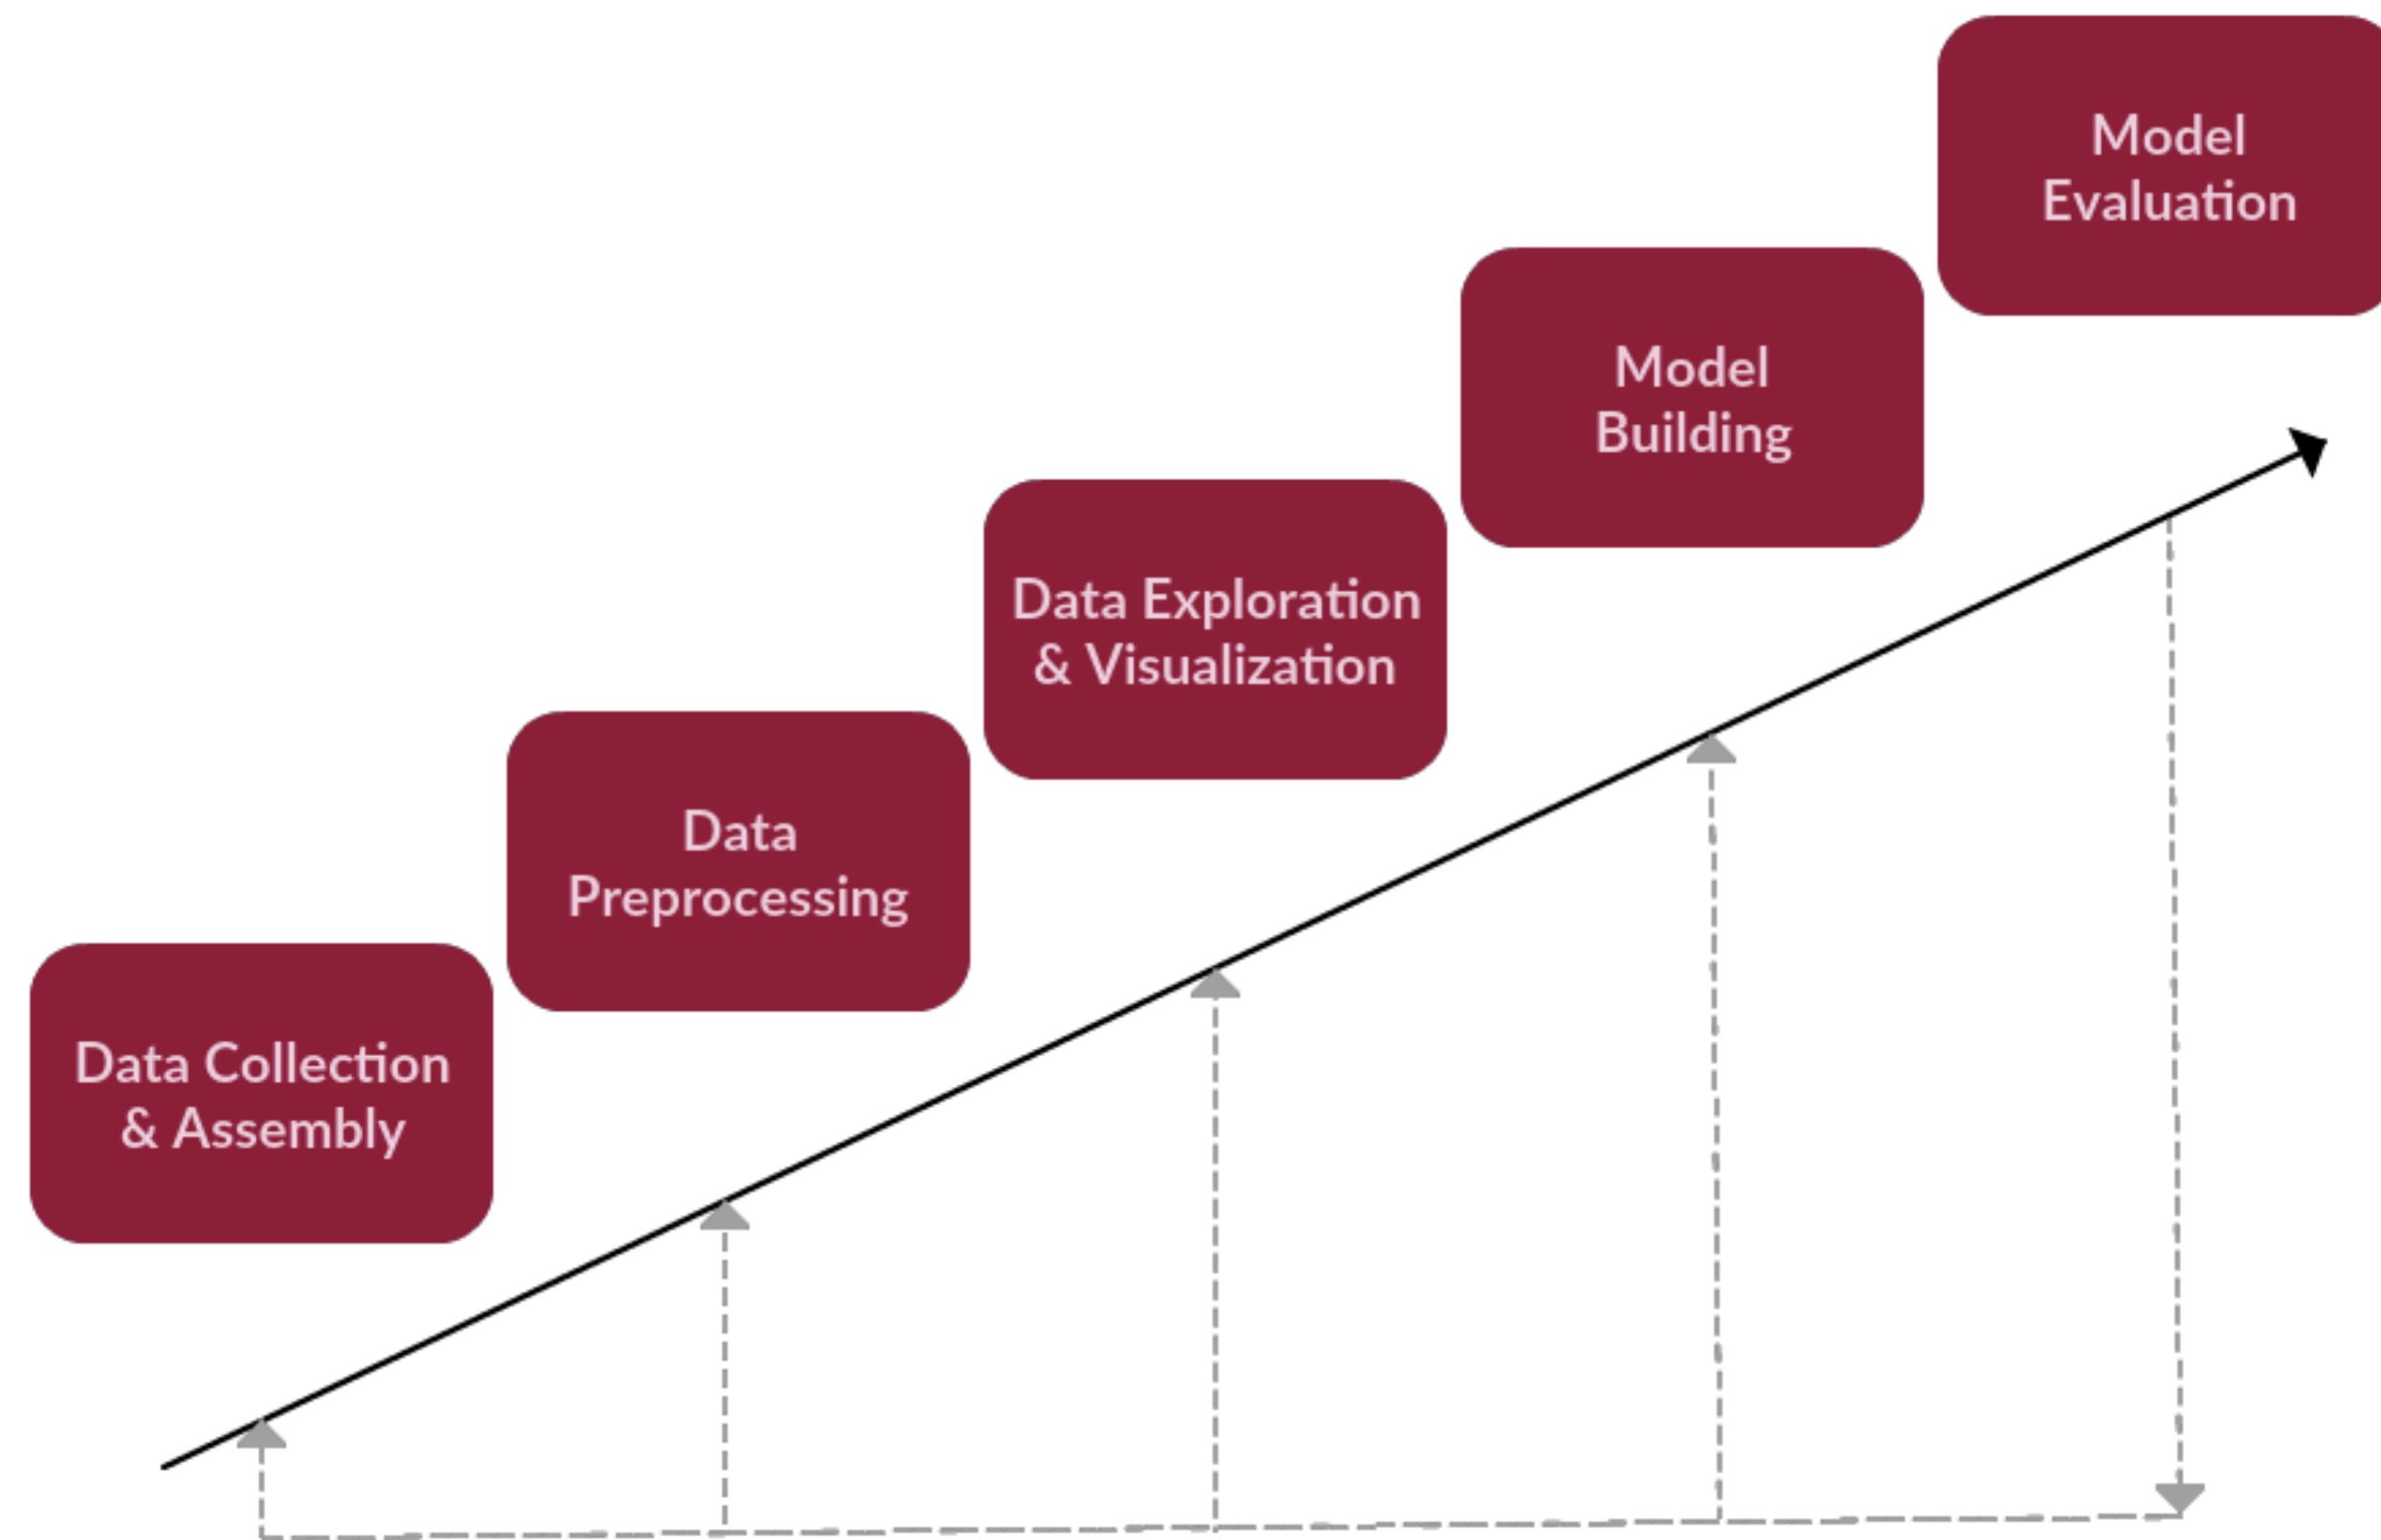
magnimind


## ***NLP Review***

# DATA TYPES



# TEXT DATA SCIENCE FRAMEWORK

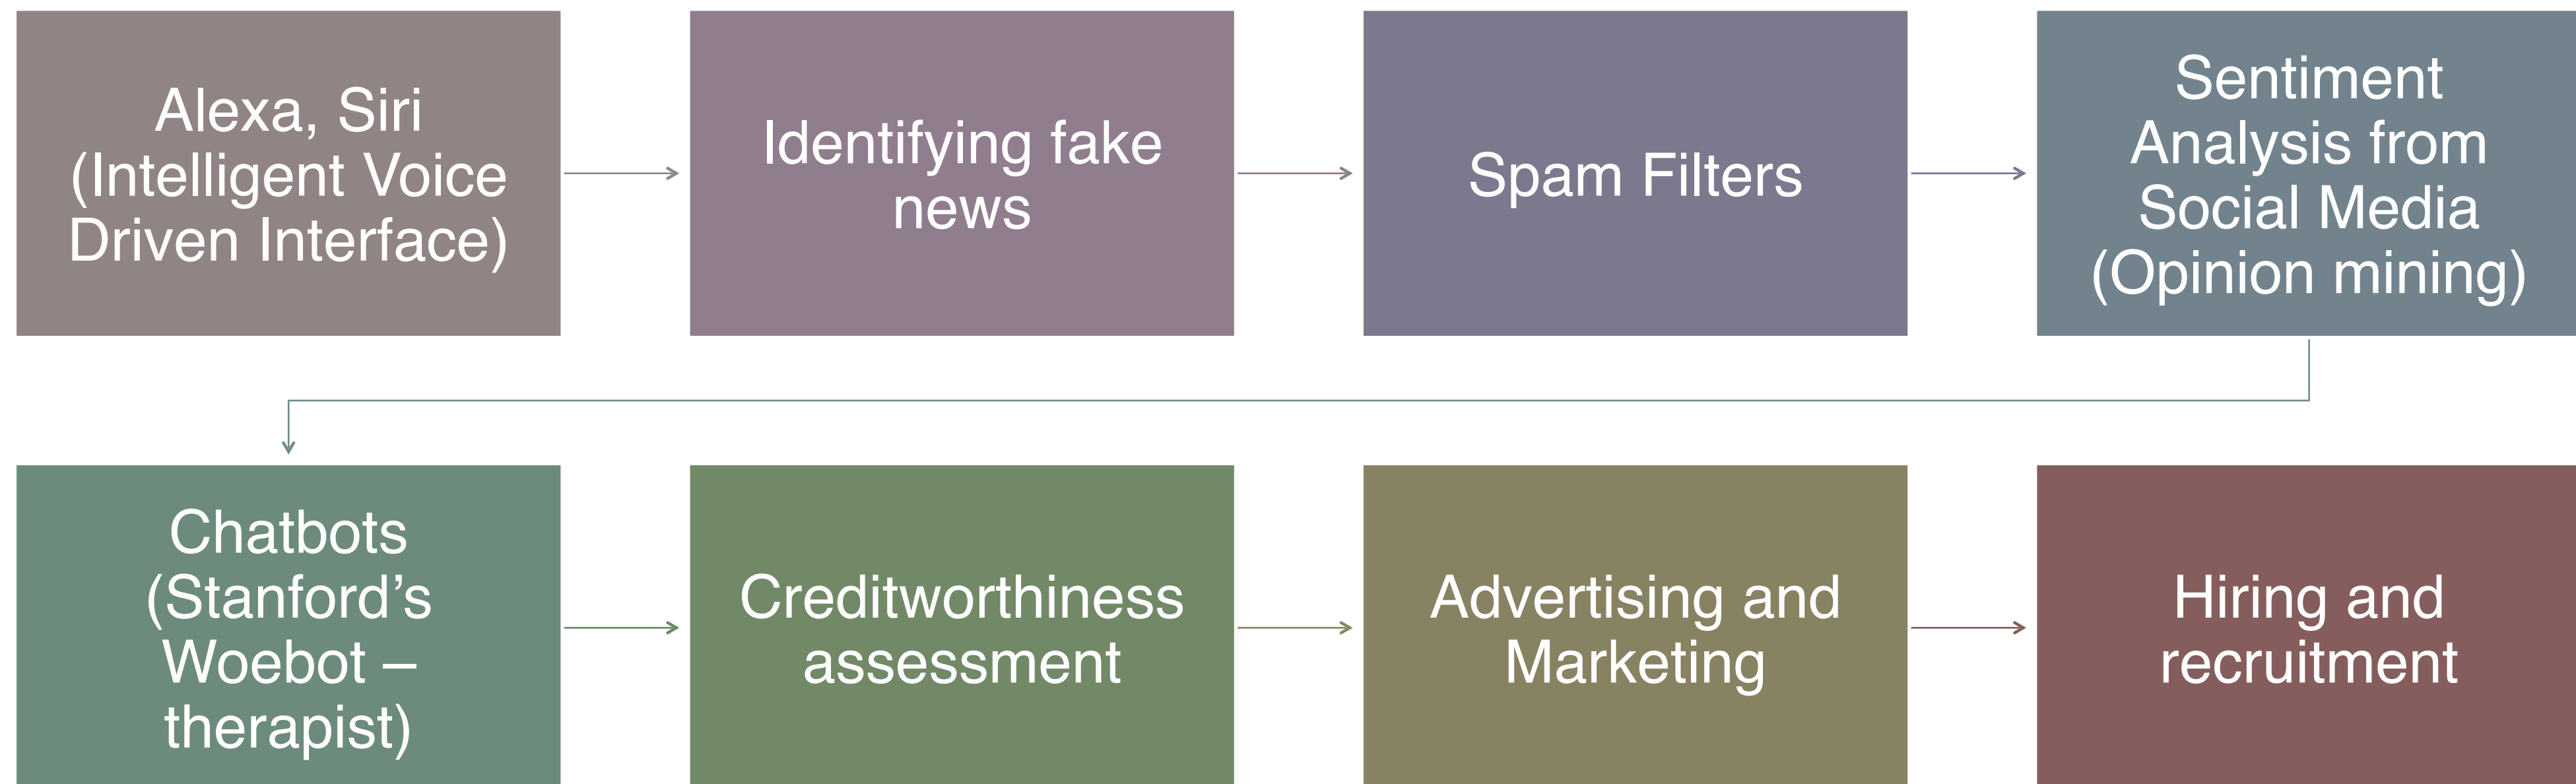




© Statista 2019


*Revenues from the natural language processing (NLP) market worldwide from 2017 to 2025 (in million U.S. dollars)*

# NLP Use Cases



# Twitter sentiment versus Gallup Poll of Consumer Confidence

Brendan O'Connor, Ramnath Balasubramanyan, Bryan R. Routledge, and Noah A. Smith. 2010. From Tweets to Polls: Linking Text Sentiment to Public Opinion Time Series. In ICWSM-2010








## Twitter sentiment:

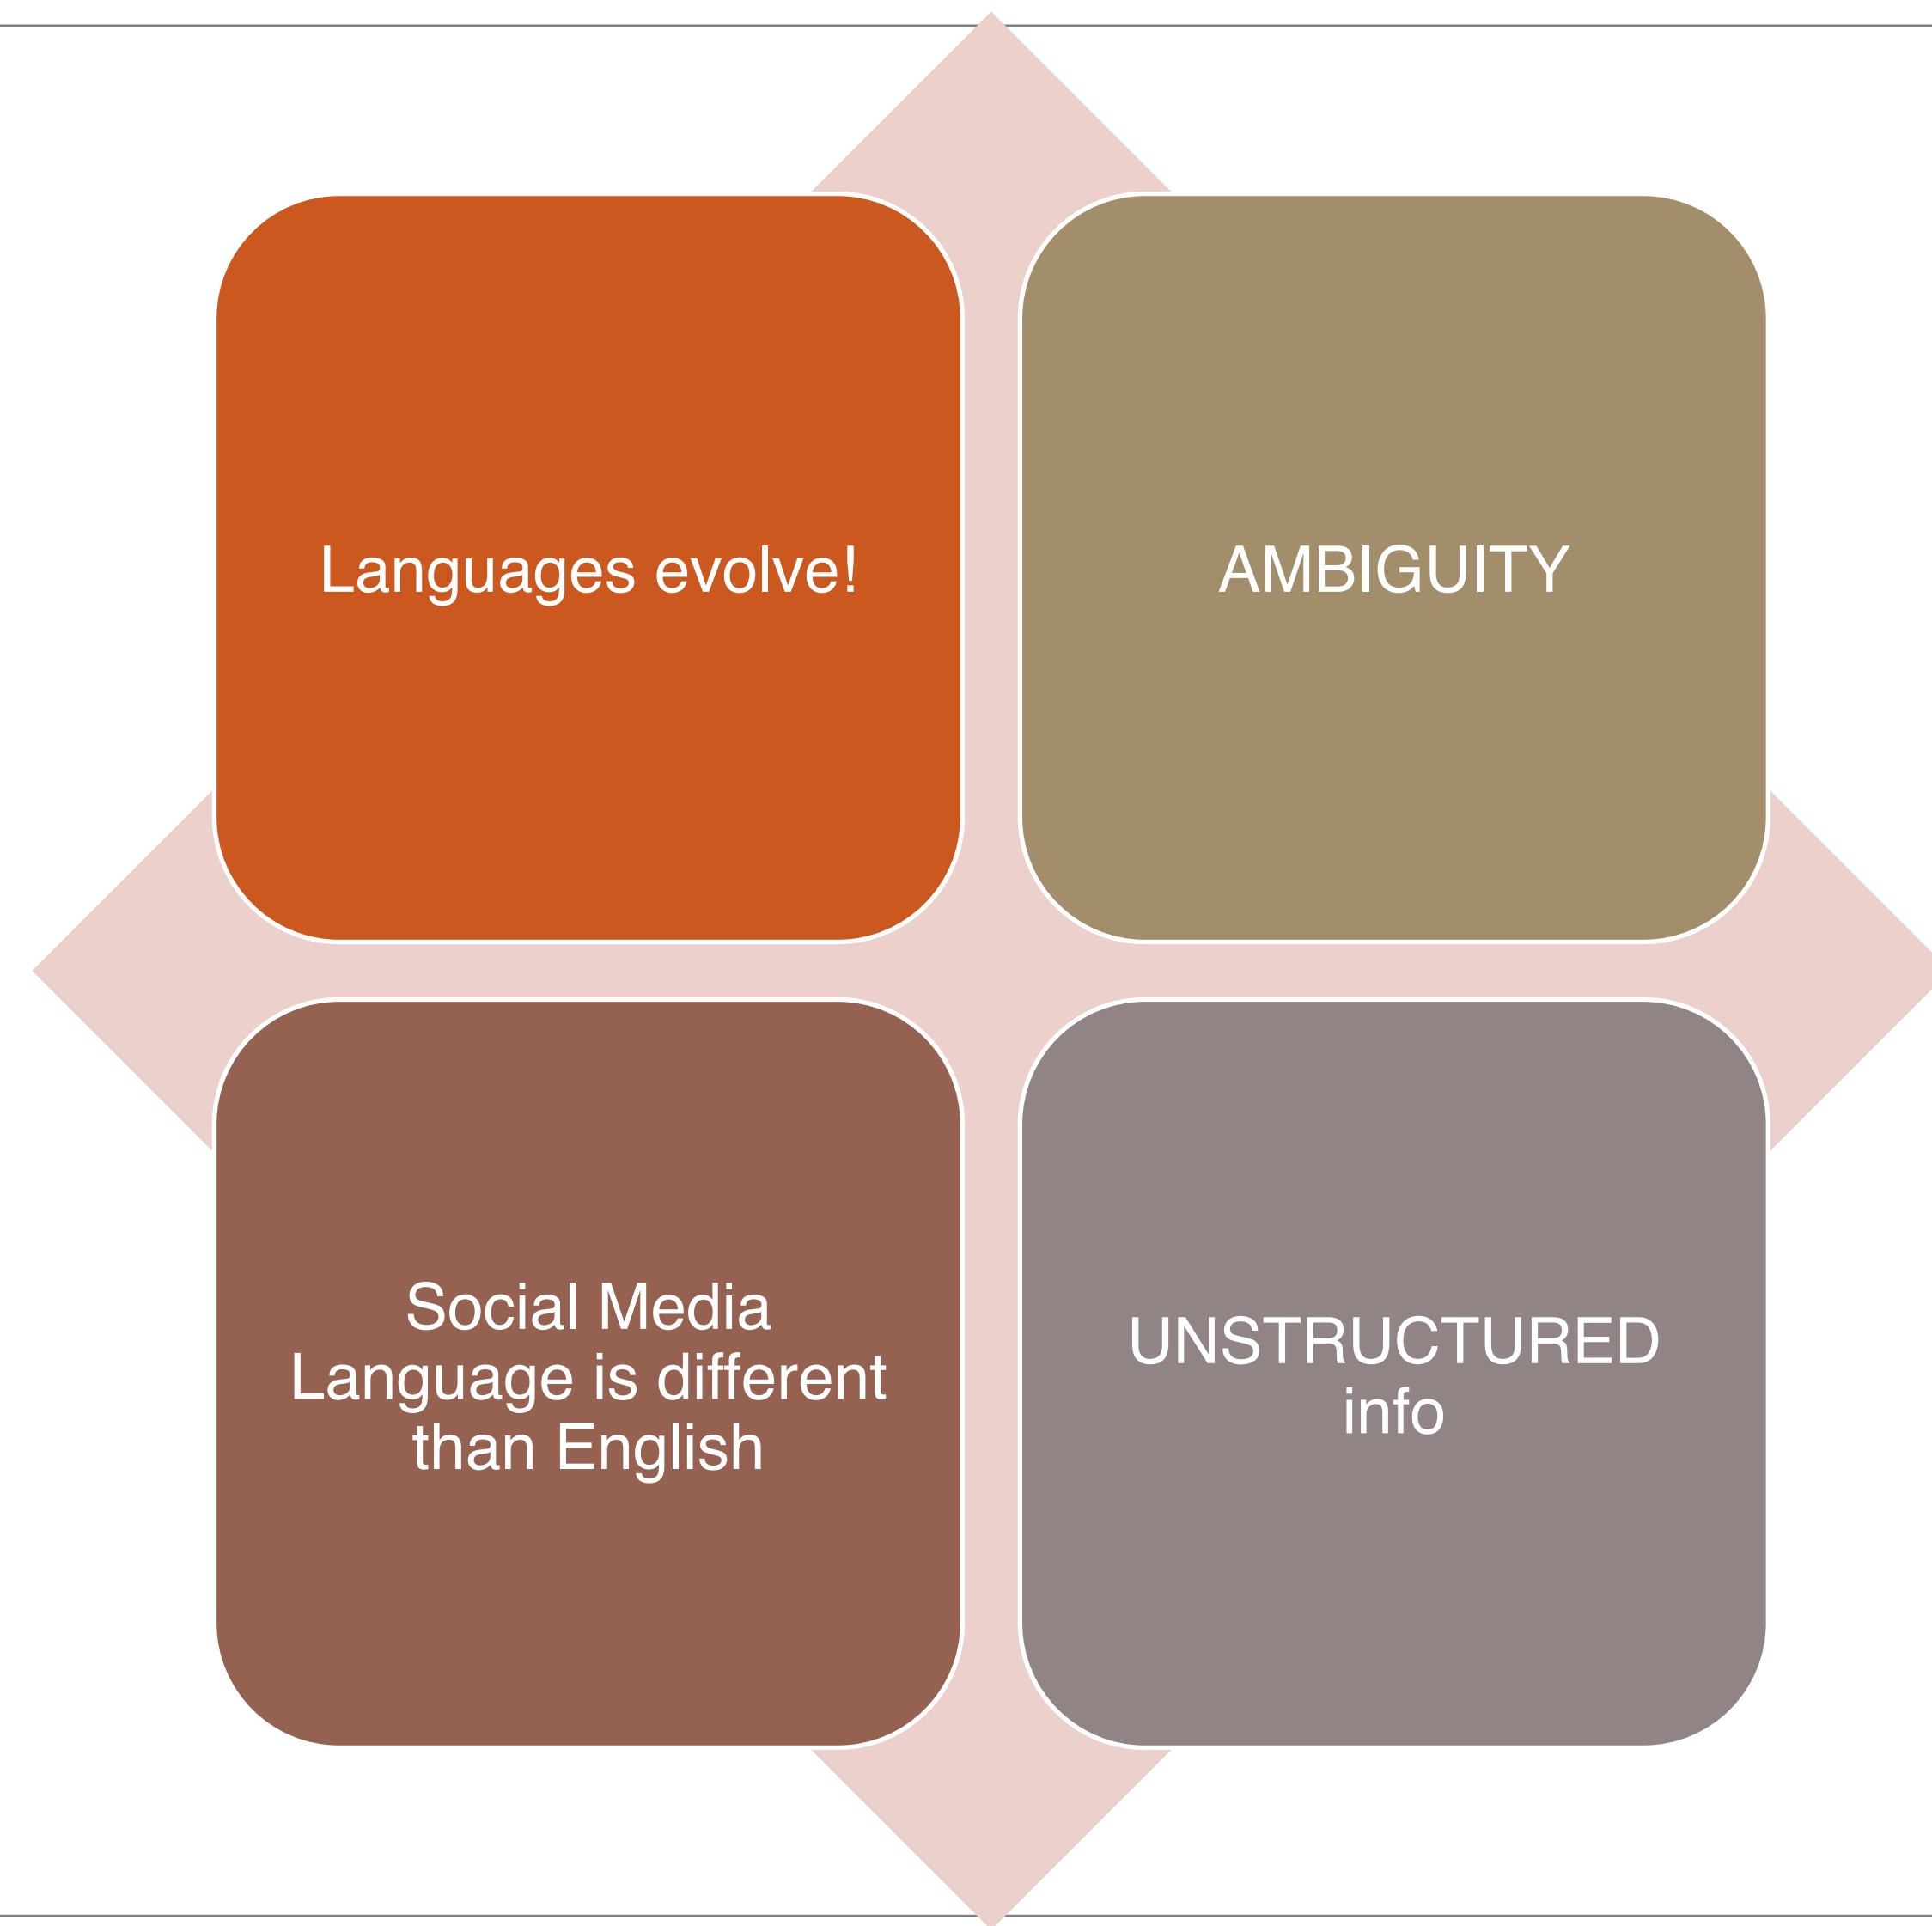
Johan Bollen, Huina Mao, Xiaojun Zeng. 2011.

Twitter mood predicts the stock market,

Journal of Computational Science 2:1, 1-8.

10.1016/j.jocs.2010.12.007.

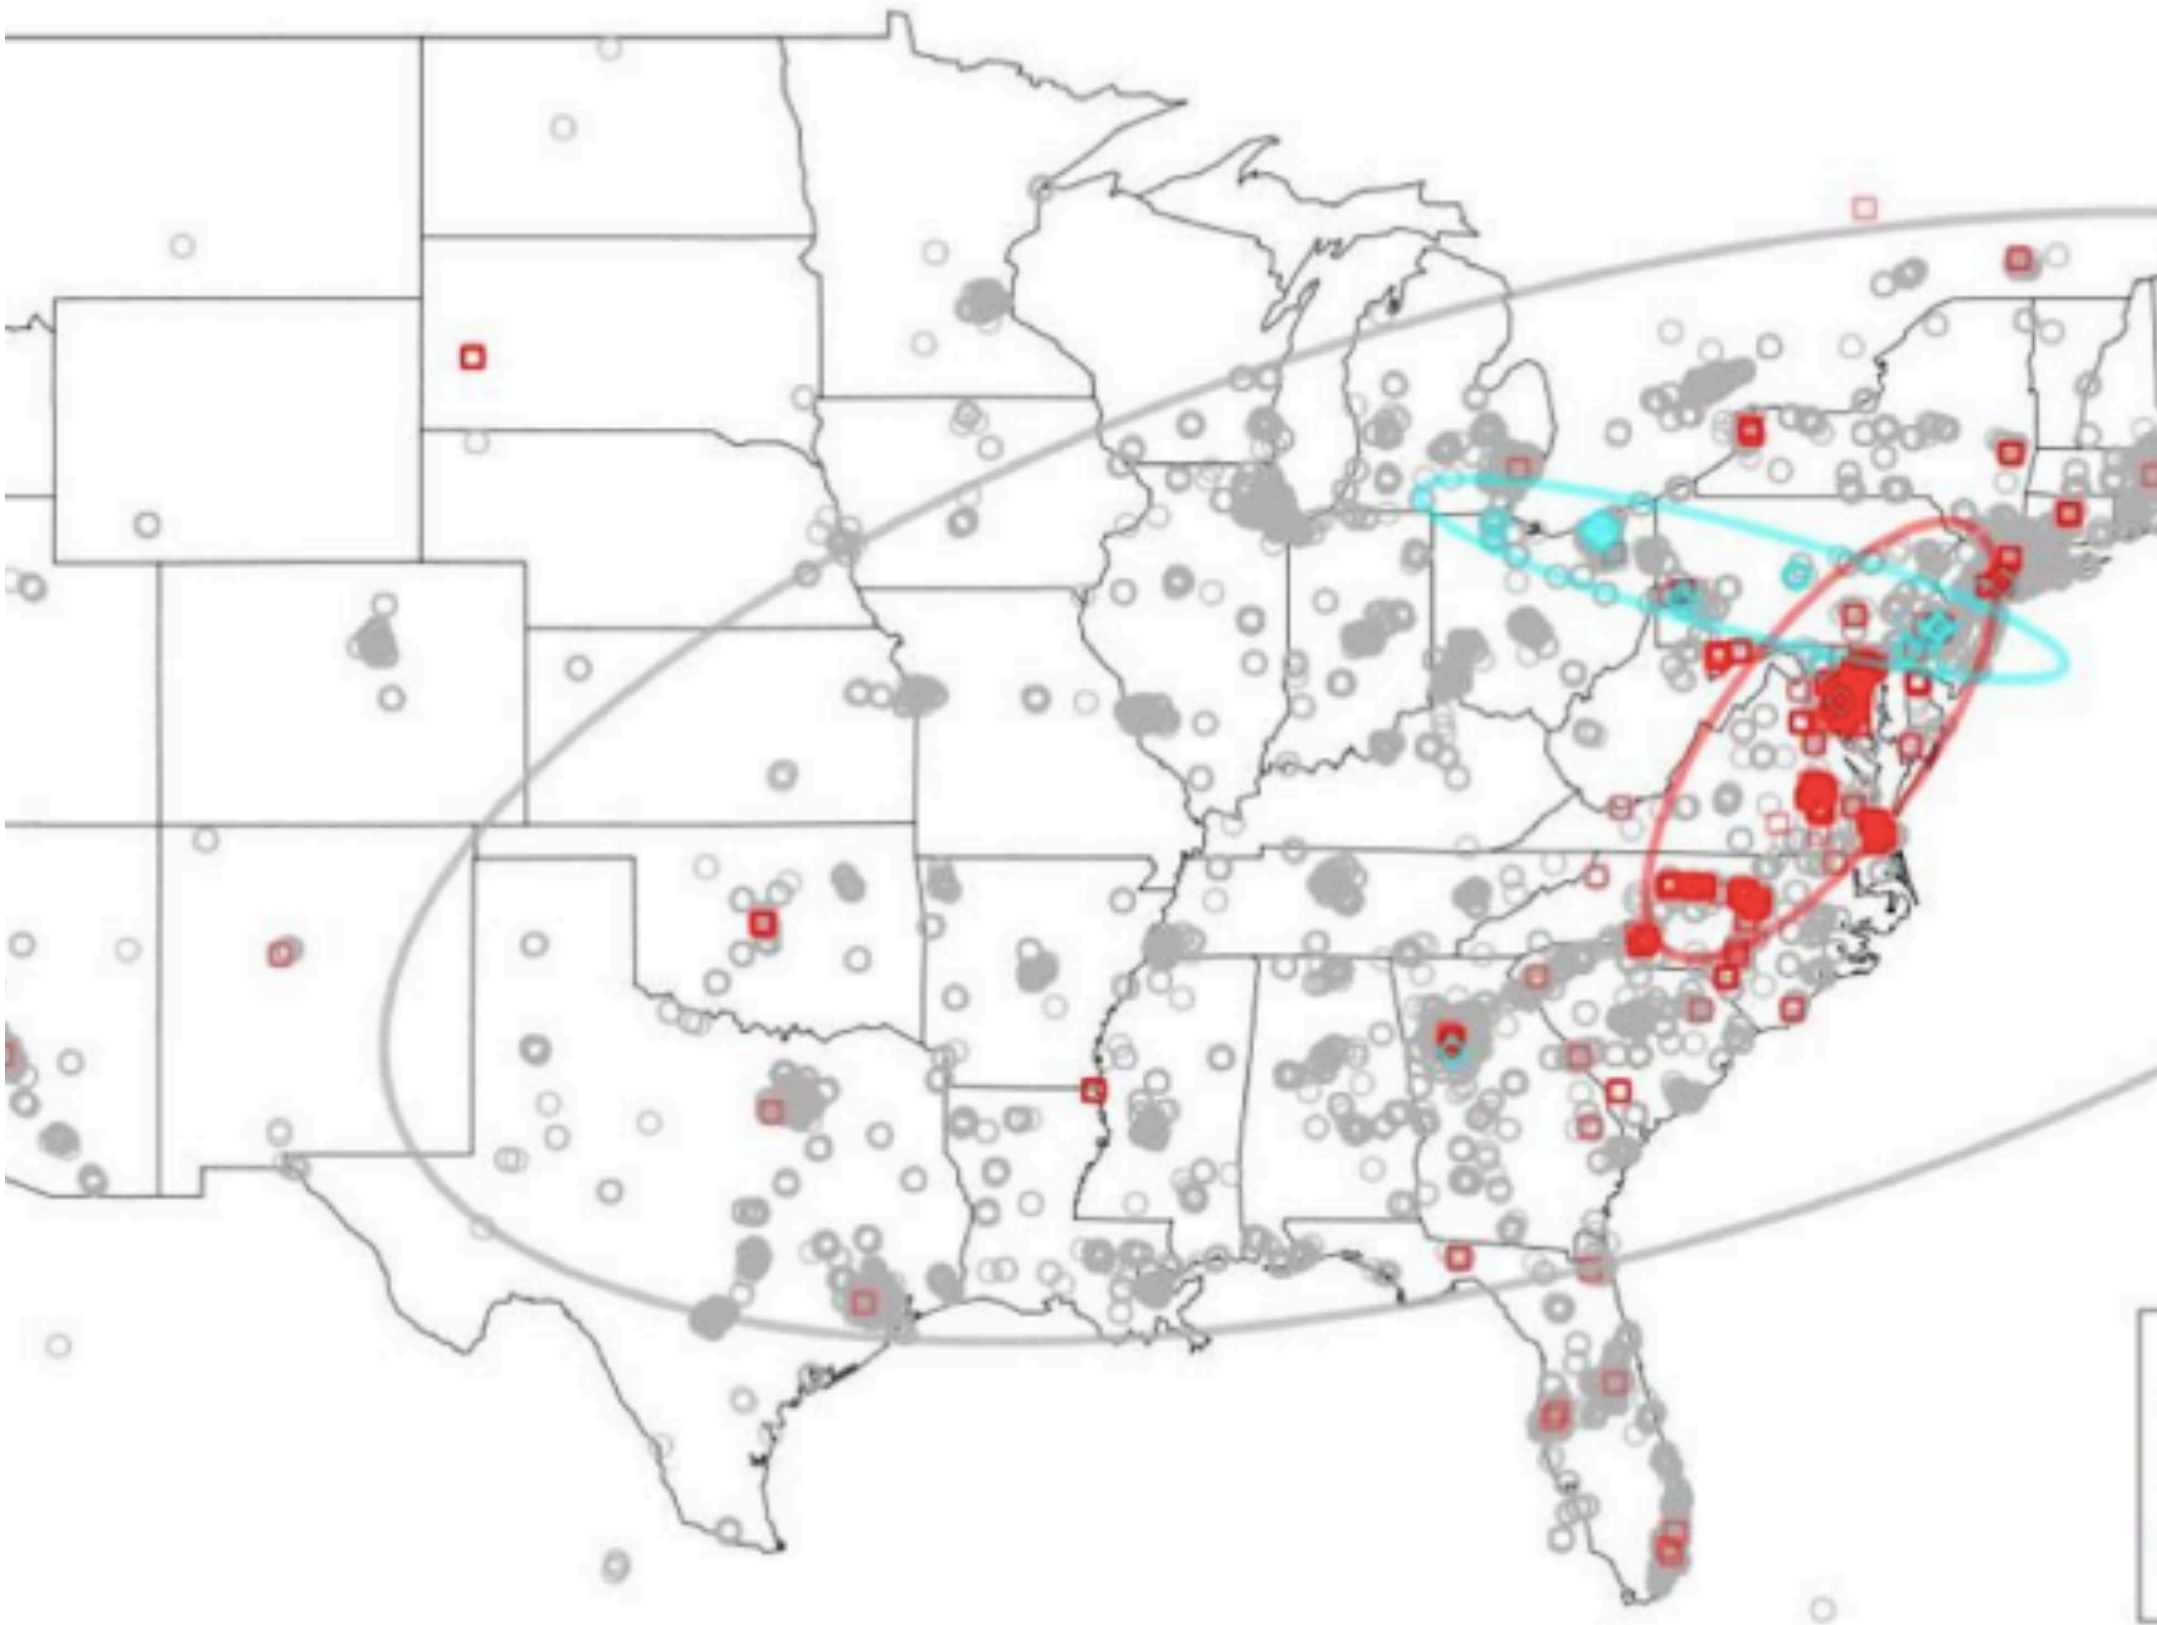








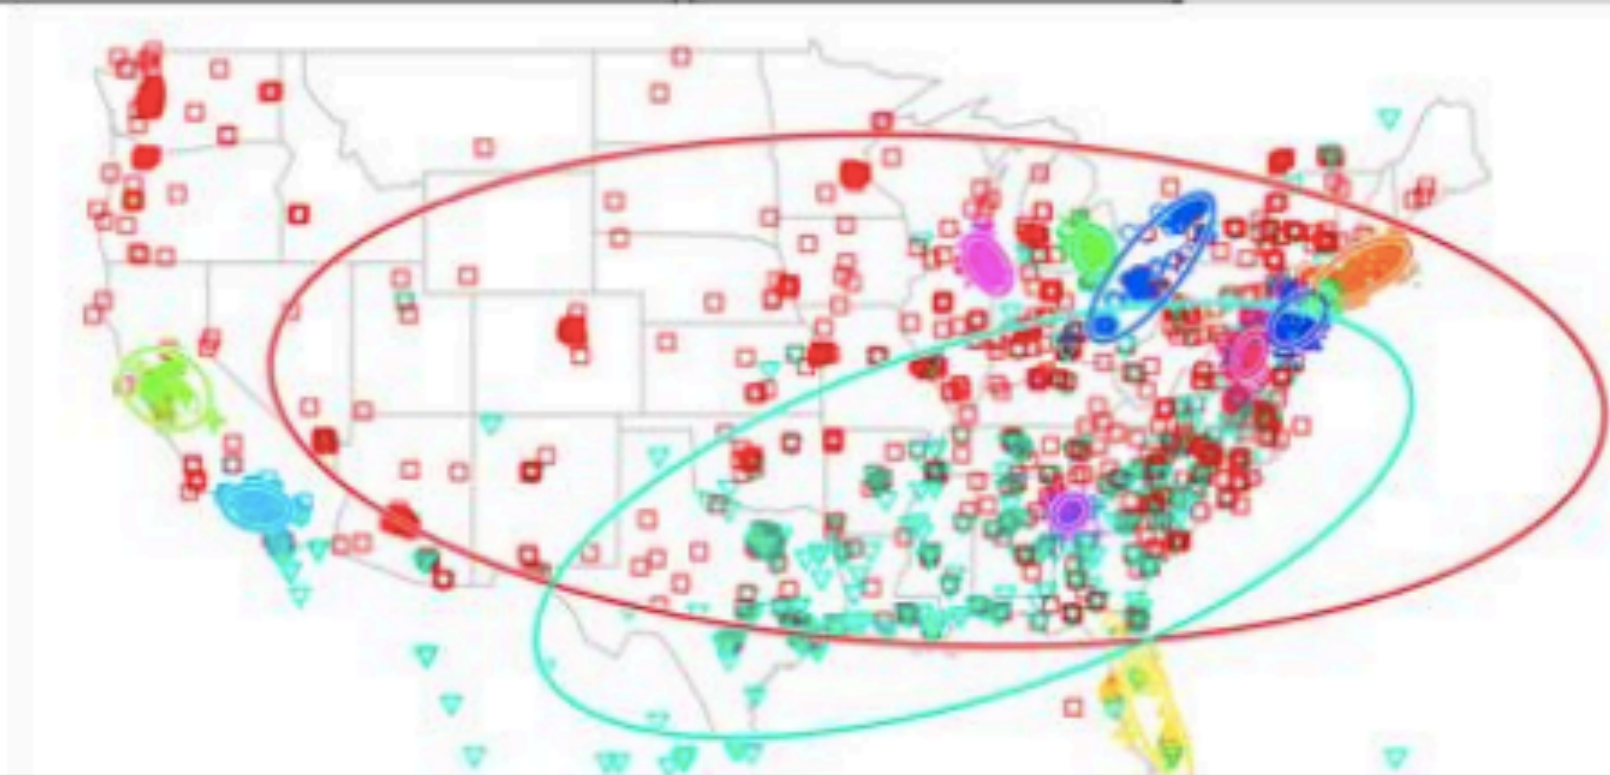
# What Makes NLP Hard?



# graphic Variation, Slang



	“basketball”	“chit chat”
	PISTONS KOBE LAKERS game DUKE NBA CAVS STUCKEY JETS KNICKS	lol smh jk yea wyd coo ima wassup somethin jp
Boston 	CELTICS victory BOSTON CHARLOTTE	<i>ese</i> exam suttin sippin
N. California 	THUNDER KINGS GIANTS pimp trees clap	hella flirt hut iono OAKLAND
New York 	NETS KNICKS	wassup nm
Los Angeles 	#KOBE #LAKERS AUSTIN	wyd coo af <i>nada</i> tacos messin fasho bomb



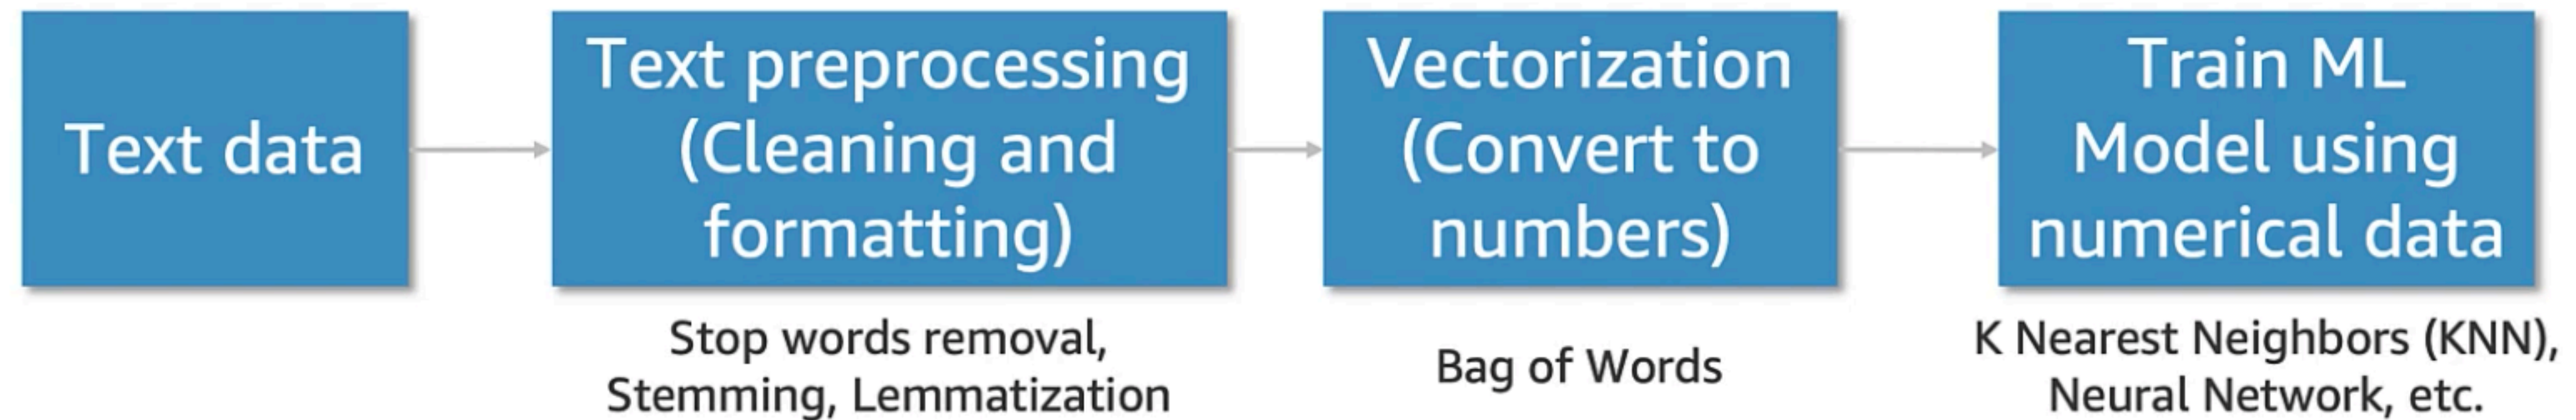
# Microtext Features

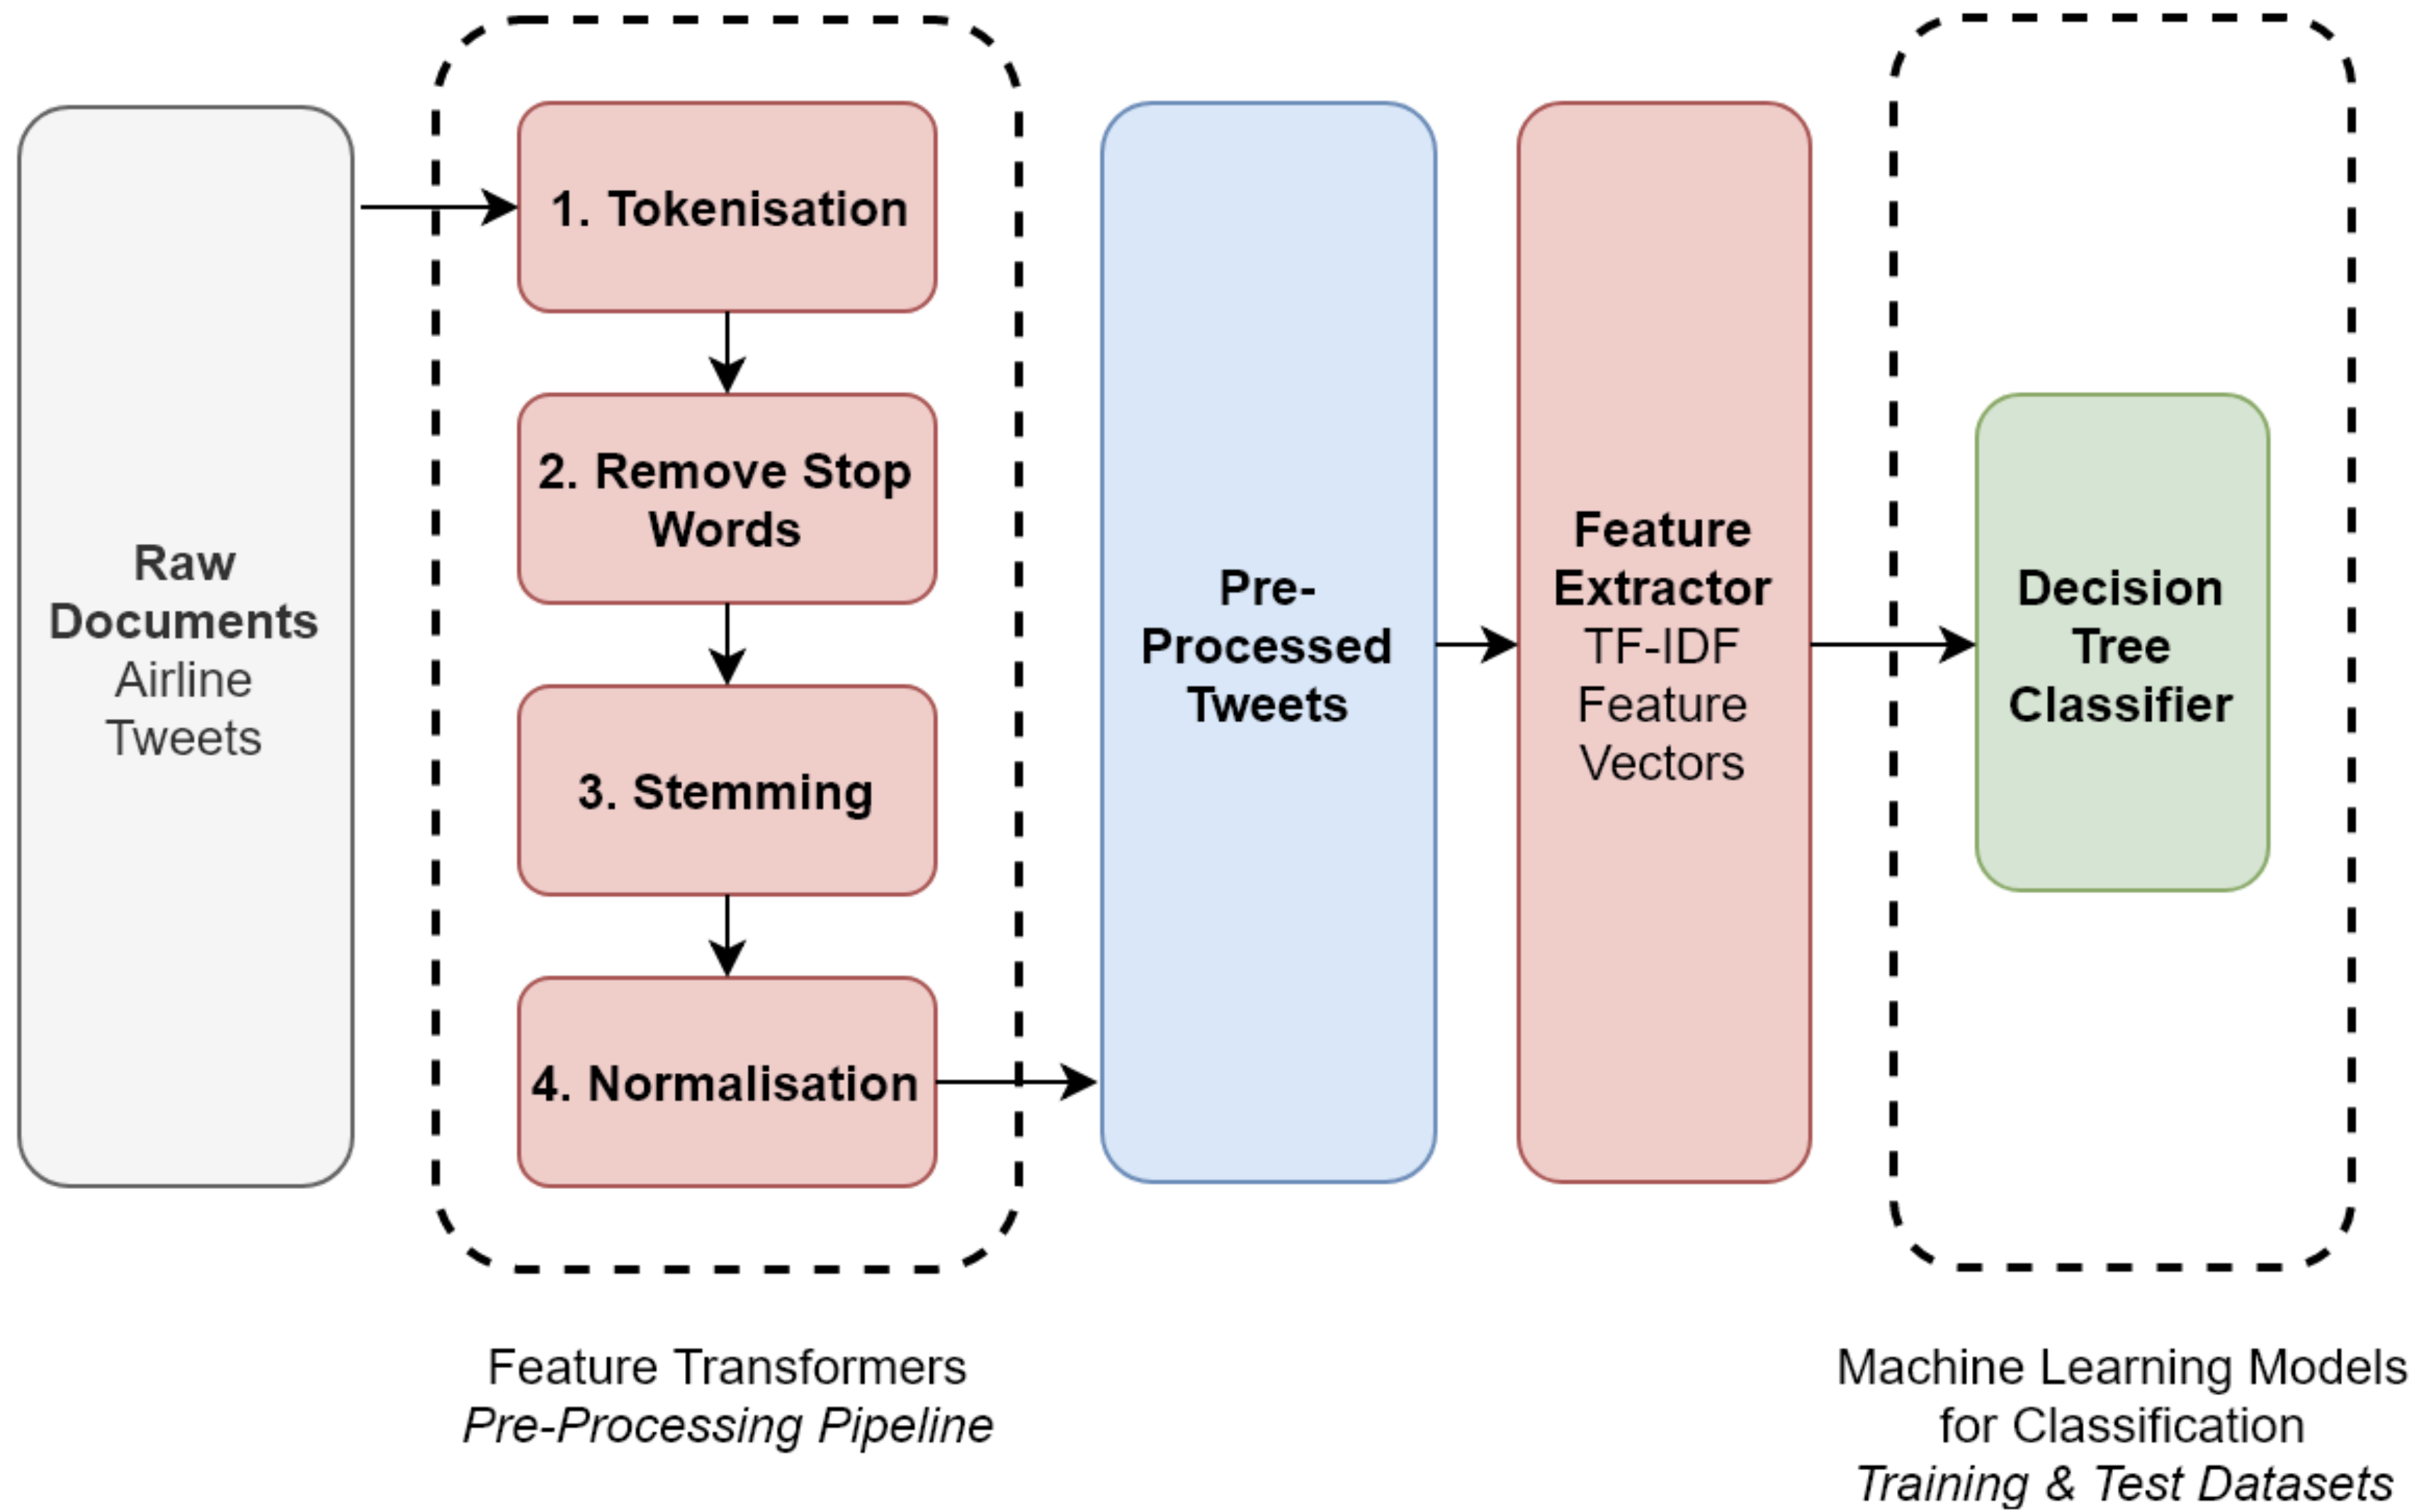
- Highly relaxed spelling
- Reliance on emoticons
- Out-of-vocabulary (OOV) words
- Phonetic spellings (b4 for before)
- Emotional emphasis (Cooooooooool)
- Popular acronyms (OTW – On the way)



# Machine Learning with Text Data

ML models need **well-defined numerical data**.






# Some NLP related words

- **Corpus:** Large collection of words or phrases - can come from different sources: documents, web sources, database
  - [Common Crawl Corpus](#): web crawl data composed of over 5 billion web pages (541 TB)
  - [Reddit Submission Corpus](#): publicly available Reddit submissions (42 GB)
  - [Wikipedia XML Data](#): complete copy of all Wikimedia wikis, in the form of wikitext source and metadata embedded in XML. (500 GB)
  - Etc.



# Some NLP Terms



**Token:** Words or phrases  
extracted from documents



# Tokenization

Splits text/document into small parts by white space and punctuation.

**Example:**

Sentence	Tokens
"I don't like eggs."	"I", "do", "n't", "like", "eggs", "."

These tokens will be used in the next steps in the pipeline.

# Stop Word Removal

**Stop words:** Some words that frequently appear in texts, but they don't contribute too much to the overall meaning.

- Common stop words: "a", "the", "so", "is", "it", "at", "in", "this", "there", "that", "my"
- **Example:**

Original sentence	Without stop words
"There is a tree near the house"	"tree near house"

- **Stemming** : chopping the affixes

compressed -> compress

compression -> compress

### **Porter's Stemmer**

- **Lemmatization** : reducing the words to their base forms

am, is, are -> be

car, cars, car's, cars' -> car

# Theory of NLP

Corpus

Document – entity/unit/object

Text segmentation

Tokenization

Word - Term

**Terms are features** of the doc

Each term has properties

normalized form of the term -> term.baseform

position(s) in the doc -> term.position(s)

frequency of the term -> **term.frequency.**

# Converting Words to Terms

- Preprocess and normalize the words  
.tolower(), stemming, lemmatization

# Feature Representation: Bag of Words

the dog is on the table

0	0	1	1	0	1	1	1
are	cat	dog	is	now	on	table	the

A single word is a one-hot encoding vector with the size of the dictionary :(



# Bag of words

- simplified and effective way to process documents by:
  - disregarding grammar (term.baseform?)
  - disregarding word order (term.position)
  - keeping only multiplicity (term.frequency)

# Bag of words

- sparse matrix
- numbers can be:
  - binary - 0/1
  - simple term frequency
  - weight - e.g. TF-IDF

Terms \ Documents	Terms				
	Investment Risk	Project Management	Software Engineering	Development	SAP
Document 1	1			1	
Document 2		1			
Document 3			3		1
Document 4		1			
Document 5			2	1	
Document 6	1			1	
...					

# Word Counts in the vector

Simple example using word counts:

	a	cat	dog	happy	is	it	my	not	old	wolf
"It is a dog."	1	0	1	0	1	1	0	0	0	0
"my cat is old"	0	1	0	0	1	0	1	0	1	0
"It is not a dog, it a is wolf."	2	0	1	0	2	2	0	1	0	1

# TF in the vector

**Term frequency (TF):** Increases the weight for **common** words in a document.

$$tf(term, doc) = \frac{\text{number of times the term occurs in the doc}}{\text{total number of terms in the doc}}$$

	a	cat	dog	is	it	my	not	old	wolf
"It is a dog."	0.25	0	0.25	0.25	0.25	0	0	0	0
"my cat is old"	0	0.25	0	0.25	0	0.25	0	0.25	0
"It is not a dog, it a is wolf."	0.22	0	0.11	0.22	0.22	0	0.11	0	0.11

# TF-IDF

term	idf
a	$\log(3/3)+1=1$
cat	$\log(3/2)+1=1.18$
dog	$\log(3/3)+1=1$
is	$\log(3/4)+1=0.87$
it	$\log(3/3)+1=1$
my	$\log(3/2)+1=1.18$
not	$\log(3/2)+1=1.18$
old	$\log(3/2)+1=1.18$
wolf	$\log(3/2)+1=1.18$

**Inverse document frequency (IDF):** Decreases the weights for **commonly** used words and **increases** weights for **rare** words in the vocabulary.

$$idf(term) = \log\left(\frac{n_{documents}}{n_{documents\ containing\ the\ term} + 1}\right) + 1$$

$$e.g.\ idf("cat") = 1.18$$

**Term Freq. Inverse Doc. Freq (TF-IDF):** Combines term frequency and inverse document frequency.

$$tf_{idf}(term, doc) = tf(term, doc) * idf(term)$$

	a	cat	dog	is	it	my	not	old	wolf
"It is a dog."	0.25	0	0.25	0.22	0.25	0	0	0	0
"my cat is old"	0	0.3	0	0.22	0	0.3	0	0.3	0
"It is not a dog, it a is wolf."	0.22	0	0.11	0.19	0.22	0	0.13	0	0.13

# Tf-Idf

---

## Term Frequency — Inverse Document Frequency

---

a technique to quantify a word in documents based on its relevancy

---

we generally compute a weight to each word which signifies the importance of the word in the document and corpus

---

This method is a widely used technique in Information Retrieval, keyword extraction, and Text Mining



- **TF-IDf** = Term Frequency (TF) \* Inverse Document Frequency (IDF)  
(0:no word in the doc, 1:the doc has only the word)
- **Tf** -> Term frequency : frequency of a word in a doc  
Tf = count of t in doc/count of the words in the doc
- **Df** -> Document frequency: count of t in the document set N  
(normalize df – divide with the number of documents)

$$\text{Idf} = N/\text{df}$$

$$\text{Idf} = \log(N/(\text{df}+1))$$

$$\text{Tf-Idf} = \text{tf} / \log(N/(\text{df}+1))$$



# Example of Tf-IDf

- Tf- Consider a document containing 100 words wherein the word *cat* appears 3 times. The term frequency (i.e., tf) for *cat* is then  $(3 / 100) = 0.03$ .
- Idf- Now, assume we have 10 million documents and the word *cat* appears in one thousand of these. Then, the inverse document frequency (i.e., idf) is calculated as  $\log(10,000,000 / 1,000) = 4$ .
- Tf-Idf: Thus, the Tf-idf weight is the product of these quantities:  $0.03 * 4 = 0.12$ .

# Basic string manipulation

- keep it simple and stupid

```
.lower(), .strip(), .split(), .join(),  
iterators, ...
```

- regexp

- not only match, but transformation, extraction (\1), backreferences etc.

- re.options, re.multiline, repl can be function:

```
def repl(m): ...  
  
re.sub("pattern", repl, "string")
```

# spaCy

- Industrial strength
- Faster than NLTK, CoreNLP, ZPar...
- Easy to install, simple
- Interoperates with Tensorflow, Keras, Scikit-Learn, Gensim

# Stanford NLP

- <http://nlp.stanford.edu/software/index.shtml>
- statistical NLP, deep learning NLP, and rule-based NLP tools for major computational linguistics problems
- famous
- Java

# Scikit-Learn

<http://scikit-learn.org/stable/index.html>

machine learning in python

## Classification

Identifying to which category an object belongs to.

**Applications:** Spam detection, Image recognition.

**Algorithms:** SVM, nearest neighbors, random forest, ... — Examples

## Regression

Predicting a continuous-valued attribute associated with an object.

**Applications:** Drug response, Stock prices.

**Algorithms:** SVR, ridge regression, Lasso, ... — Examples

## Clustering

Automatic grouping of similar objects into sets.

**Applications:** Customer segmentation, Grouping experiment outcomes

**Algorithms:** k-Means, spectral clustering, mean-shift, ... — Examples

## Dimensionality reduction

Reducing the number of random variables to consider.

**Applications:** Visualization, Increased efficiency

**Algorithms:** PCA, feature selection, non-negative matrix factorization. — Examples

## Model selection

Comparing, validating and choosing parameters and models.

**Goal:** Improved accuracy via parameter tuning

**Modules:** grid search, cross validation, metrics. — Examples

## Preprocessing

Feature extraction and normalization.

**Application:** Transforming input data such as text for use with machine learning algorithms.

**Modules:** preprocessing, feature extraction. — Examples