

Spam filtering

Many of the e-mail services dedicate a special folder in their inboxes for the e-mails that are categorized as spams. The underlying mechanism that works behind the scenes is a text classification model which takes the content of the e-mail as input and returns a probability regarding this e-mail being a spam. This area of NLP is called **text classification** and we'll walk you through some applications of text categorization in the remainder of code.

```
In [1]: import pandas as pd
```

Load and simplify the dataset

```
In [37]: df = pd.read_csv(r'spam.csv', encoding='ISO-8859-1')

df.rename(columns = {'v1': 'class_label', 'v2': 'message'}, inplace = True)
df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis = 1, inplace = True)

# notice your class_labels

df
```

```
Out[37]:
```

	class_label	message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...
...
5567	spam	This is the 2nd time we have tried 2 contact u...
5568	ham	Will i_b going to esplanade fr home?
5569	ham	Pity, * was in mood for that. So...any other s...
5570	ham	The guy did some bitching but I acted like i'd...
5571	ham	Rofl. Its true to its name

5572 rows x 2 columns

```
In [38]: # exploring the dataset. Have a look at your class counts.  
# This will give you an idea whether you have an imbalanced  
# dataset problem  
df['class_label'].value_counts()
```

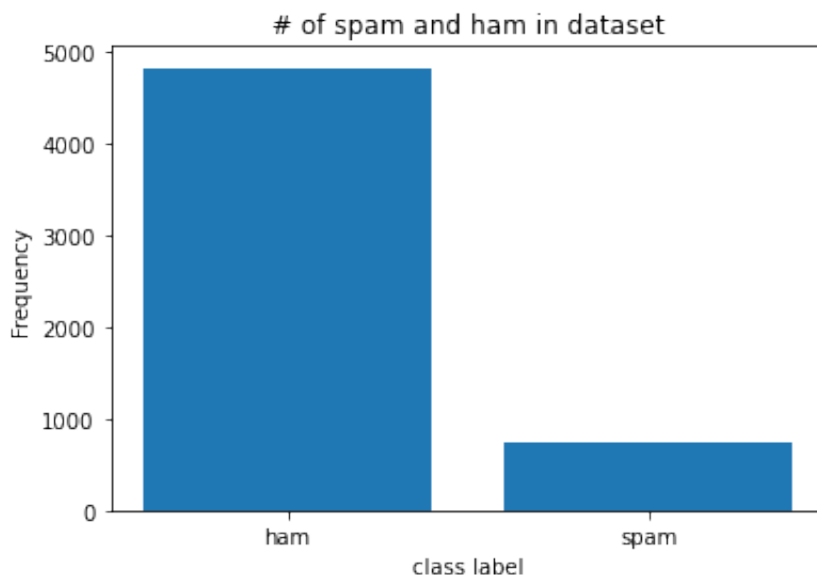
```
Out[38]: ham      4825  
spam       747  
Name: class_label, dtype: int64
```

```
In [39]: #exploratory data analysis  
#bar chart for initial data  
import matplotlib.pyplot as ab  
import numpy as np  
  
labels = ['ham', 'spam']  
counts = [4825, 747]  
ypos = np.arange(len(labels)) #converting text labels to numeric value, 0 and 1  
ypos
```

```
Out[39]: array([0, 1])
```

```
In [40]: ab.xticks(ypos, labels)  
ab.xlabel("class label")  
ab.ylabel("Frequency")  
ab.title("# of spam and ham in dataset")  
ab.bar(ypos, counts)
```

```
Out[40]: <BarContainer object of 2 artists>
```



```
In [41]: df
```

Out[41]:

	class_label	message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...
...
5567	spam	This is the 2nd time we have tried 2 contact u...
5568	ham	Will i_b going to esplanade fr home?
5569	ham	Pity, * was in mood for that. So...any other s...
5570	ham	The guy did some bitching but I acted like i'd...
5571	ham	Rofl. Its true to its name

5572 rows × 2 columns

```
In [42]: #separate both classes
df_ham = df[df.class_label=='ham']
df_spam = df[df.class_label=='spam']

#convert the dataframe to a list. This is how you will print the wordcloud.
ham_list=df_ham['message'].tolist()
spam_list= df_spam['message'].tolist()
#spam_list
```

```
In [43]: filtered_spam = ("").join(spam_list) #convert the list into a string of spam
filtered_spam = filtered_spam.lower()

filtered_ham = ("").join(ham_list) #convert the list into a string of ham
filtered_ham = filtered_ham.lower()
#filtered_spam
```

```
In [44]: import os
from wordcloud import WordCloud
from PIL import Image
comment_mask = np.array(Image.open("comment.png"))
```

```
# Create and generate a word cloud image:
wordcloud = WordCloud(max_font_size = 160, margin=0, mask = comment_mask, background_color="black")

# Display the generated image

ab.figure(figsize=[20,10])
ab.imshow(wordcloud, interpolation='bilinear')
ab.axis("off")

ab.margins(x=0, y=0)

ab.show()
```



```
ab.show()
```



```
# since class labels are strings, convert them to numbers: 1 for spam
df['class_label'] = df['class_label'].apply(lambda x: 1 if x == 'spam' else 0)
```

```
In [48]: from sklearn.model_selection import train_test_split
# 30% test, 70% train split
# split the dataset
x_train, x_test, y_train, y_test = train_test_split(df['message'], df['class_

print('rows in test set: ' + str(x_test.shape))
print('rows in train set: ' + str(x_train.shape))
type(x_train)

rows in test set: (1672,)
rows in train set: (3900,)

Out[48]: pandas.core.series.Series
```

tfidf Vectorizer cell

```
In [49]: from sklearn.feature_extraction.text import TfidfVectorizer

# Fitting of the CountVectorizer consists of tokenization of the training dat
# and building the vocabulary
# Transforming the CountVectorizer creates the bag-of-words representation
# of the train data.
# the bow is stored in a SciPy sparse matrix that only stores the nonzero
# entries.
# to look at the actual content of the sparse matrix, convert it to dense
# array using numpy.toarray() method

# vectorize email text into tfidf matrix
# TfidfVectorizer converts collection of raw documents to a matrix of TF-IDF
# It's equivalent to CountVectorizer followed by TfidfTransformer.
list = x_train.tolist()
vectorizer = TfidfVectorizer(
    input=list, # input is actual text
    lowercase=True, # convert to lower case before tokenizing
    stop_words='english' # remove stop words
)
features_train_transformed = vectorizer.fit_transform(list) #gives tf idf vec
features_test_transformed = vectorizer.transform(x_test) #gives tf idf vecto
# after transform, boW with tfidf values are calculated.
```

```
In [50]: from sklearn.naive_bayes import MultinomialNB

# train a classifier
classifier = MultinomialNB()
classifier.fit(features_train_transformed, y_train)
```

```
Out[50]: MultinomialNB()
```

```
In [51]: # review the classifier accuracy
print("classifier accuracy {:.2f}%".format(classifier.score(features_test_tra
```

classifier accuracy 96.41%

```
In [52]: # accuracy looks good but it is not a good indicator
# for imbalanced datasets.
```

```
In [53]: labels = classifier.predict(features_test_transformed)
from sklearn.metrics import f1_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report

actual = y_test.tolist()
predicted = labels
results = confusion_matrix(actual, predicted)

print('Confusion Matrix :')
print(results)
print ('Accuracy Score :',accuracy_score(actual, predicted))
print ('Report : ')
print (classification_report(actual, predicted) )

score_2 = f1_score(actual, predicted, average = 'binary')
print('F-Measure: %.3f' % score_2)
```

Confusion Matrix :

```
[[1434    0]
 [  60  178]]
```

Accuracy Score : 0.9641148325358851


Report :

	precision	recall	f1-score	support
0	0.96	1.00	0.98	1434
1	1.00	0.75	0.86	238
accuracy			0.96	1672
macro avg	0.98	0.87	0.92	1672
weighted avg	0.97	0.96	0.96	1672

F-Measure: 0.856

```
In [54]: #heatmap for confusion matrix
import seaborn as sns
group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']
group_counts = ["{0:0.0f}".format(value) for value in
                 results.flatten()]
group_percentages = ["{0:.2%}".format(value) for value in
                     results.flatten()/np.sum(results)]
labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
          zip(group_names,group_counts,group_percentages)]
labels = np.asarray(labels).reshape(2,2)
sns.heatmap(results, annot=labels, fmt='', cmap='Reds')
```

Out[54]: <AxesSubplot:>



count vectorizer cell

```
In [55]: list = x_train.tolist()
from sklearn.feature_extraction.text import CountVectorizer
cv = CountVectorizer(strip_accents = ascii) #initialize countVectorizer.
word_count_vector=cv.fit_transform(list) #get the document-term matrix.
print(word_count_vector.shape) #see the rows and columns of this matrix

(3900, 7127)
```

```
In [ ]: #cv.vocabulary_
```

```
In [57]: x_test_vector = cv.transform(x_test)
```

```
In [58]: # train a classifier
classifier = MultinomialNB()
classifier.fit(word_count_vector, y_train)
```

Out[58]: MultinomialNB()

```
In [59]: # review the classifier accuracy
print("classifier accuracy {:.2f}%".format(classifier.score(x_test_vector, y_test)))

classifier accuracy 98.44%
```



```
In [60]: labels = classifier.predict(x_test_vector)

actual = y_test
predicted = labels
results = confusion_matrix(actual, predicted)

print('Confusion Matrix :')
print(results)
print ('Accuracy Score :',accuracy_score(actual, predicted))
print ('Report : ')
print (classification_report(actual, predicted) )
score_2 = f1_score(actual, predicted, average = 'binary')
print('F-Measure: %.3f' % score_2)
```

Confusion Matrix :

```
[[1427   7]
 [  19 219]]
```

Accuracy Score : 0.9844497607655502

Report :


	precision	recall	f1-score	support
0	0.99	1.00	0.99	1434
1	0.97	0.92	0.94	238
accuracy			0.98	1672
macro avg	0.98	0.96	0.97	1672
weighted avg	0.98	0.98	0.98	1672

F-Measure: 0.944

```
In [61]: #heatmap for confusion matrix

group_names = ['True Neg','False Pos','False Neg','True Pos']
group_counts = ["{0:0.0f}".format(value) for value in
                results.flatten()]
group_percentages = ["{0:.2%}".format(value) for value in
                    results.flatten()/np.sum(results)]
labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
          zip(group_names,group_counts,group_percentages)]
labels = np.asarray(labels).reshape(2,2)
sns.heatmap(results, annot=labels, fmt='', cmap='Greens')
```

Out[61]: <AxesSubplot:>



```
In [62]: #Implementing cross validation on both the models

#THIS CELL USES CROSS VALIDATION FOR COUNT VECTORIZER

#%pip install scikit-learn

#%conda upgrade scikit-learn

from sklearn.model_selection import cross_val_score

X=df["message"].tolist()
Y=df["class_label"].tolist()
df

classifier = MultinomialNB()

cv = CountVectorizer(strip_accents = ascii) #initialize countVectorizer.
word_count_vector=cv.fit_transform(X) #get the document-term matrix.
#print(word_count_vector.shape) #see the rows and columns of this matrix

scores=cross_val_score(classifier,word_count_vector,Y,cv=10,scoring='accuracy')

print(scores)

print(scores.mean())

[0.98387097 0.98028674 0.98025135 0.98384201 0.98025135 0.98025135
 0.98204668 0.98384201 0.97127469 0.98922801]
0.9815145138768233
```

```
In [63]: from sklearn.model_selection import cross_val_predict


y_pred = cross_val_predict(classifier, word_count_vector, Y, cv=10)
conf_mat = confusion_matrix(Y, y_pred)
conf_mat
```

```
Out[63]: array([[4758,   67],
               [   36,  711]])
```

```
In [64]: #heatmap for confusion matrix

group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']
group_counts = ["{0:0.0f}".format(value) for value in
                conf_mat.flatten()]
group_percentages = ["{0:.2%}".format(value) for value in
                    conf_mat.flatten()/np.sum(conf_mat)]
labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
          zip(group_names, group_counts, group_percentages)]
labels = np.asarray(labels).reshape(2,2)
sns.heatmap(conf_mat, annot=labels, fmt='', cmap='terrain')
```

```
Out[64]: <AxesSubplot:>
```



```
In [65]: vectorizer = TfidfVectorizer(
    input= list , # input is actual text
    lowercase=True, # convert to lower case before tokenizing
    stop_words='english' # remove stop words
)
features_train_transformed = vectorizer.fit_transform(X) #gives tf idf vector
#features_test_transformed = vectorizer.transform(x_test) #gives tf idf vect

classifier = MultinomialNB()

scores=cross_val_score(classifier,features_train_transformed,Y,cv=10,scoring=

print(scores)

print(scores.mean())

[0.98387097 0.97849462 0.96947935 0.97845601 0.96947935 0.96768402
 0.97127469 0.97307002 0.97307002 0.98025135]
0.9745130402887975
```

```
In [66]: y_pred = cross_val_predict(classifier,features_train_transformed, Y, cv=10)
conf_mat = confusion_matrix(Y, y_pred)
conf_mat
```

```
Out[66]: array([[4824,    1],
               [ 141,   606]])
```

```
In [67]: #heatmap for confusion matrix

group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']
group_counts = ["{0:0.0f}".format(value) for value in
                conf_mat.flatten()]
group_percentages = ["{0:.2%}".format(value) for value in
                     conf_mat.flatten()/np.sum(conf_mat)]
labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
          zip(group_names,group_counts,group_percentages)]
labels = np.asarray(labels).reshape(2,2)
sns.heatmap(conf_mat, annot=labels, fmt='', cmap='twilight_shifted_r')
```

Out[67]: <AxesSubplot:>



In []:

In []: